

Analysis

Question 1:

Part 1: are the benchmark timings for StringStrand consistent with the explanation below that the time to execute cutAndSplice is $O(b^2S)$?

Note that the value of b is half the number of calls to append since each cut (except the first) is modeled by two calls of append in the method cutAndSplice -- see the code. This means that b^2 will be constant in the benchmark, but S will vary.

Answer 1 Part 1:

Knowing that b^2 is constant we can analyze how the runtime changes with respect to S changing. S is the size of the splicee, and we see very clearly when the splicee is hundreds of thousands in size that when it doubles the runtime doubles which is keeping consistent with our expectations. This is because Strings are immutable objects which cannot be changed, so to when splicing a new String (more memory) must be used and then populated every time a break occurs. This leads to the most inefficient function out of the three tested.

Answer 1 Part 2:

After modifying the Benchmark program to triple the mySource, we expect b to be tripled which we see from the first few lines of the new Benchmark. This should result in our new runtime of $O((3b)^2s)$. Below are those times which match with our expectations of what would happen when increasing b in an $O(b^2s)$ runtime function.

Results from DNABenchmark using mySource = mySource+mySource+mySource;

Class	splicee	recomb	time	appends
StringStrand:	256	14,401,413	7.402	3870
StringStrand:	512	14,896,773	8.397	3870
StringStrand:	1,024	15,887,493	6.960	3870
StringStrand:	2,048	17,868,933	9.516	3870
StringStrand:	4,096	21,831,813	11.768	3870

Results from DNABenchmark using StringStrand on ecoli.txt

dna length = 4,639,221
cutting at enzyme gaattc

Class	splicee	recomb	time	appends
StringStrand:	256	4,800,471	0.384	1290
StringStrand:	512	4,965,591	0.356	1290
StringStrand:	1,024	5,295,831	0.359	1290
StringStrand:	2,048	5,956,311	0.418	1290

StringStrand:	4,096	7,277,271	0.526	1290
StringStrand:	8,192	9,919,191	0.892	1290
StringStrand:	16,384	15,203,031	1.319	1290
StringStrand:	32,768	25,770,711	2.629	1290
StringStrand:	65,536	46,906,071	5.173	1290
StringStrand:	131,072	89,176,791	9.774	1290
StringStrand:	262,144	173,718,231	19.623	1290
StringStrand:	524,288	342,801,111	38.633	1290
StringStrand:	1,048,576	680,966,871	74.557	1290

Exception in thread "main" java.lang.OutOfMemoryError: Java heap space

at java.base/jdk.internal.misc.Unsafe.allocateUninitializedArray(Unsafe.java:1269)
at

java.base/java.lang.invoke.StringConcatFactory\$MethodHandleInlineCopyStrategy.newArray(StringConcatFactory.java:1633)

at java.base/java.lang.invoke.DirectMethodHandle\$Holder.invokeStatic(DirectMethodHandle\$Holder)

at java.base/java.lang.invoke.LambdaForm\$MH/0x00000000800149840.invoke(LambdaForm\$MH)

at java.base/java.lang.invoke.Invokers\$Holder.linkToTargetMethod(Invokers\$Holder)

at StringStrand.append(StringStrand.java:70)

at IDnaStrand.cutAndSplice(IDnaStrand.java:37)

at DNABenchmark.lambda\$strandSpliceBenchmark\$0(DNABenchmark.java:76)

at DNABenchmark\$\$Lambda\$163/0x0000000080025f840.run(Unknown Source)

at java.base/java.lang.Thread.run(Thread.java:834)

at DNABenchmark.strandSpliceBenchmark(DNABenchmark.java:79)

at DNABenchmark.main(DNABenchmark.java:112)

Process finished with exit code 1

Question 2: are the benchmark timings for StringBuilderStrand consistent with the explanation below that the time to execute cutAndSplice is $O(bs)$?

Answer 2 Part 1:

We can see below the output of StringBuilderStrand which we expect to have the same runtimes as in question 2 with StringStrand since b is held constant as can be seen by each append remaining at 1290. The reason why this class is more efficient than StringStrand is because it makes use of StringBuilder which uses a char array to populate the the new strand which since it is primitive does not require as much memory as a String because primitives are mutable.

Answer 2 Part 2:

Below are the first few runs of the Benchmark with b tripled. We can see a clear difference between these runs and the runs when b was 1290. As was expected from theory, the majority of the runs grew larger because the runtime $O(bs)$ would have changed to $O(3bs)$. However as can be seen from closer analysis of the runtimes, the time did not triple exactly which can be assumed to be caused by the state computer system itself during the trials.

Results from DNABenchmark using `mySource = mySource+mySource+mySource;`

Class	splicee	recomb	time appends
StringBuilderStrand:	256	14,401,413	0.038 3870
StringBuilderStrand:	512	14,896,773	0.039 3870
StringBuilderStrand:	1,024	15,887,493	0.041 3870

StringBuilderStrand:	2,048	17,868,933	0.043	3870
StringBuilderStrand:	4,096	21,831,813	0.047	3870
StringBuilderStrand:	8,192	29,757,573	0.041	3870
StringBuilderStrand:	16,384	45,609,093	0.056	3870
StringBuilderStrand:	32,768	77,312,133	0.107	3870
StringBuilderStrand:	65,536	140,718,213	0.099	3870
StringBuilderStrand:	131,072	267,530,373	0.275	3870
StringBuilderStrand:	262,144	521,154,693	0.442	3870
StringBuilderStrand:	524,288	1,028,403,333	0.738	3870

Results from DNABenchmark using StringBuilderStrand on ecoli.txt

dna length = 4,639,221
cutting at enzyme gaattc

```
-----
Class          splicee      recomb    time appends
-----
StringBuilderStrand:      256      4,800,471 0.023      1290
StringBuilderStrand:      512      4,965,591 0.018      1290
StringBuilderStrand:     1,024      5,295,831 0.007      1290
StringBuilderStrand:     2,048      5,956,311 0.007      1290
StringBuilderStrand:     4,096      7,277,271 0.008      1290
StringBuilderStrand:     8,192      9,919,191 0.009      1290
StringBuilderStrand:    16,384     15,203,031 0.011      1290
StringBuilderStrand:    32,768     25,770,711 0.022      1290
StringBuilderStrand:    65,536     46,906,071 0.029      1290
StringBuilderStrand:   131,072     89,176,791 0.077      1290
StringBuilderStrand:   262,144    173,718,231 0.101      1290
StringBuilderStrand:   524,288    342,801,111 0.255      1290
StringBuilderStrand:  1,048,576    680,966,871 0.415      1290
```

```
Exception in thread "main" java.lang.OutOfMemoryError: Java heap space
  at java.base/java.util.Arrays.copyOf(Arrays.java:3745)
  at java.base/java.lang.AbstractStringBuilder.ensureCapacityInternal(AbstractStringBuilder.java:172)
  at java.base/java.lang.AbstractStringBuilder.append(AbstractStringBuilder.java:538)
  at java.base/java.lang.StringBuilder.append(StringBuilder.java:174)
  at StringBuilderStrand.append(StringBuilderStrand.java:70)
  at IDnaStrand.cutAndSplice(IDnaStrand.java:41)
  at DNABenchmark.strandSpliceBenchmark(DNABenchmark.java:67)
  at DNABenchmark.main(DNABenchmark.java:112)
```

Process finished with exit code 1

Question 3: Explain why the time for LinkStrand does not change much at all over all the runs in the benchmark program. Explain why you think memory is exhausted at the specific strand size you see in your timings -- as compared to exhaustion for String and StringBuilder.

Answer 3

The reason that LinkStrand does not change much during all the runs is because we are taking advantage of how nodes work. With each node having info of a sequence of chars, the only work the computer has to do is to change the pointer for the info of the node

which runs at $O(1)$, constant time. Another benefit is that the splicee does not have to be stored again in memory. Instead the pointers of the nodes can point to the same splicee.

Results from DNABenchmark using LinkStrand on ecoli.txt

dna length = 4,639,221
cutting at enzyme gaattc

```
-----
Class          splicee      recomb    time appends
-----
LinkStrand:    256          4,800,471 0.044      1290
LinkStrand:    512          4,965,591 0.029      1290
LinkStrand:    1,024        5,295,831 0.011      1290
LinkStrand:    2,048        5,956,311 0.006      1290
LinkStrand:    4,096        7,277,271 0.006      1290
LinkStrand:    8,192        9,919,191 0.007      1290
LinkStrand:    16,384       15,203,031 0.006      1290
LinkStrand:    32,768       25,770,711 0.013      1290
LinkStrand:    65,536       46,906,071 0.006      1290
LinkStrand:    131,072      89,176,791 0.006      1290
LinkStrand:    262,144     173,718,231 0.013      1290
LinkStrand:    524,288     342,801,111 0.006      1290
LinkStrand:    1,048,576   680,966,871 0.007      1290
LinkStrand:    2,097,152   1,357,298,391 0.013      1290
LinkStrand:    4,194,304   2,709,961,431 0.007      1290
LinkStrand:    8,388,608   5,415,287,511 0.010      1290
LinkStrand:    16,777,216 10,825,939,671 0.012      1290
LinkStrand:    33,554,432 21,647,243,991 0.004      1290
LinkStrand:    67,108,864 43,289,852,631 0.015      1290
LinkStrand:    134,217,728 86,575,069,911 0.013      1290
LinkStrand:    268,435,456 173,145,504,471 0.010      1290
LinkStrand:    536,870,912 346,286,373,591 0.010      1290
LinkStrand:    1,073,741,824 692,568,111,831 0.010      1290
```

Exception in thread "main" java.lang.OutOfMemoryError: Java heap space

```
at java.base/java.util.Arrays.copyOf(Arrays.java:3745)
at java.base/java.lang.AbstractStringBuilder.ensureCapacityInternal(AbstractStringBuilder.java:172)
at java.base/java.lang.AbstractStringBuilder.append(AbstractStringBuilder.java:538)
at java.base/java.lang.StringBuilder.append(StringBuilder.java:174)
at DNABenchmark.main(DNABenchmark.java:109)
```

Process finished with exit code 1